Answer Anyway: Prompt Pressure Elicits Fabricated Compliance in Large Language Models

Ari Holtzman and NeuriCo

Abstract

We study whether large language models exploit an "answer anyway" loophole: when a request is impossible or the model is uncertain, does pressure to answer push the model toward compliant-looking but epistemically wrong output? Prior work has measured deception, hallucination, truthfulness, and abstention in separate benchmark traditions, but has not directly tested whether these behaviors share a common answer-pressure dynamic. We run a controlled cross-benchmark study on SIMPLEQA, TRUTHFULQA, and impossible false-presupposition prompts from HALOGEN. We evaluate two API models, GPT-4.1 and GPT-5-MINI, under three prompt conditions: a neutral baseline, an abstention-friendly instruction, and a forced-answer instruction. Across 960 model generations, the clearest effect appears on impossible prompts. When abstention is forbidden, both models attempt an answer on 100% of impossible HALOGEN prompts, and most of those attempts are fabricated lists (96.7% for GPT-4.1 and 91.7% for GPT-5-MINI). On factual benchmarks, the effect is weaker and model-dependent. GPT-5-MINI shows a large abstention-versus-bluffing tradeoff: on SIMPLEQA, its attempt rate rises from 8% to 100% and its incorrect rate rises by 82 points under forced answering. GPT-4.1, by contrast, tends to answer regardless of instruction. These results show that loophole-like behavior is real, especially when prompts reward surface compliance over epistemic honesty.

1 Introduction

Large language models are often evaluated on whether they can produce an answer. In deployment, that objective can be wrong. For research assistants, search tools, tutoring systems, and agent pipelines, a confident but fabricated answer can be more harmful than an explicit "I don't know." Our main question is simple: when a model is uncertain, or when a request is impossible to satisfy, does prompt pressure to answer push it toward compliant-looking output instead of honest abstention?

This question matters because many real systems implicitly reward visible action. A downstream component may treat any non-empty response as progress, even if the response should have been an abstention. If models learn to satisfy the surface form of a request while ignoring the underlying epistemic constraint, then the failure is not just hallucination in the narrow sense. It is a broader answer-pressure failure in which the model exploits an easy loophole: answering anyway.

Existing work provides many pieces of this picture, but not a unified test. Research on strategic deception shows that models can preserve undesirable behavior under training pressure [Hubinger et al., 2024]. Truthfulness and factuality benchmarks measure plausible falsehoods and calibrated answering [Lin et al., 2021, Wei et al., 2024]. Hallucination benchmarks study fabricated content and verifier design [Li et al., 2023, Hu et al., 2024, Ravichander et al., 2025]. AbstentionBench shows that frontier reasoning models still fail badly on unanswerable questions [Kirichenko et al., 2025]. However, these lines of work usually evaluate one failure mode at a time. They do not directly ask whether the same answer-pressure dynamic connects factual guessing, abstention failure, and fabricated compliance on impossible requests.

We address that gap with a controlled cross-benchmark study. We hold question sets fixed and vary only the instruction pressure around answering. We compare a neutral BASELINE prompt, an ABSTAIN_ALLOWED prompt that explicitly permits "I don't know," and a FORCED_ANSWER prompt that forbids abstention. We test GPT-4.1 and GPT-5-MINI on SIMPLEQA, TRUTHFULQA, and false-presupposition prompts from HALOGEN, where the impossible subset gives a clean setting in which abstention is the correct behavior.

The main result is sharp. On impossible HALOGEN prompts, FORCED_ANSWER raises the loophole-attempt rate from 28.3% to 100% for GPT-4.1 and from 31.7% to 100% for GPT-5-MINI. Fabricated attempts rise by 75.0 and 61.7 percentage points, respectively. On factual benchmarks, the pattern depends strongly on the model. On SIMPLEQA, GPT-5-MINI moves from an 8% attempt rate under ABSTAIN_ALLOWED to 100% under FORCED_ANSWER, with an 82-point increase in incorrect answers. GPT-4.1 changes little because it already tends to answer even when abstention is encouraged.

Our contributions are:

- We define LOOPHOLEUSE as attempted compliance when abstention is epistemically correct, which lets us compare impossible requests and factual uncertainty within one framework.
- We conduct a controlled study across three benchmark families and three prompt conditions, totaling 960 model generations and 600 grader calls.
- We show that impossible structured prompts produce the strongest loophole effect, with both tested models switching to universal attempted compliance under FORCED_ANSWER.
- We identify a strong model difference on factual QA: GPT-5-MINI exhibits a large abstention-versus-bluffing tradeoff, while GPT-4.1 mostly answers regardless of instruction.

The rest of the paper is organized as follows. Section 2 situates the study in prior work on deception, truthfulness, hallucination, and abstention. Section 3 describes the benchmark slices, prompt conditions, grading pipeline, and statistical tests. Section 4 reports the main quantitative findings, and Section 5 discusses interpretation, limitations, and broader implications.

2 Related Work

Strategic deception. The strongest prior evidence for deliberate model misbehavior comes from trigger-conditioned settings. Sleeper Agents shows that models can preserve deceptive behavior through standard safety training and can even learn to hide that behavior better under adversarial training [Hubinger et al., 2024]. That work studies latent backdoors and persistent deception. Our setting is different. We do not claim strategic deception; instead, we measure an observable inference-time behavior in which the model produces compliant-looking output when abstention would be more appropriate.

Truthfulness and calibrated factual answering. TRUTHFULQA measures whether models reproduce common human falsehoods rather than truthful answers [Lin et al., 2021]. SIMPLEQA asks a complementary question: does a model "know what it knows," and can it leave a question unanswered when needed [Wei et al., 2024]. These benchmarks are directly relevant because they separate correct answers from plausible but wrong answers and, in the case of SIMPLEQA, explicitly allow non-attempts. We build on them by keeping the question set fixed while manipulating answer pressure, which makes the abstention-versus-bluffing tradeoff directly measurable.

Hallucination benchmarks and checkers. Hallucination has been studied through benchmark construction and verifier design. HaluEval provides a large benchmark for hallucination recognition across multiple tasks [Li et al., 2023]. RefChecker introduces claim-triplet verification for fine-grained hallucination detection [Hu et al., 2024]. HALOGEN goes further by releasing prompts, verifiers, and an error taxonomy across multiple domains [Ravichander et al., 2025]. We use HALOGEN's false-presupposition prompts because they offer a clean impossible-request setting where abstention is not only acceptable but correct.

Abstention under uncertainty. AbstentionBench argues that abstention remains an unsolved reliability problem, especially for reasoning-tuned models on unanswerable questions [Kirichenko et al., 2025]. Our work complements that line by experimentally varying prompt incentives. Instead of only measuring whether models abstain when they should, we test whether removing the abstention option causes a predictable shift into pseudo-compliance or fabrication.

Taken together, these lines of work suggest that deception, hallucination, truthfulness failure, and abstention failure are related but usually evaluated in isolation. Our contribution is to study them through a shared operational lens: whether answer pressure makes models choose compliant output over epistemically honest behavior.

3 Methodology

Task framing. We operationalize LOOPHOLEUSE as attempted compliance when abstention would be epistemically correct. This definition covers two cases. On factual benchmarks, abstention is appropriate when the model is not confident it knows the answer. On impossible false-presupposition prompts, abstention is the only correct behavior because the request cannot be satisfied exactly.

Models and prompt conditions. We evaluate two API models: GPT-4.1 and GPT-5-MINI. Each model is run under three prompt conditions: BASELINE, which asks for a brief direct answer; ABSTAIN_ALLOWED, which permits the model to say "I don't know" if unsure; and FORCED_ANSWER, which explicitly forbids abstention and asks for the best possible answer even under uncertainty. All generations use seed 42. The main experiment contains 960 model calls in total, and factual grading adds 600 calls to GPT-4.1-MINI.

Benchmarks. We use three benchmark families chosen to cover different forms of epistemic pressure. SIMPLEQA contains 50 short factual questions. TRUTHFULQA contains 50 misconception-heavy questions. HALOGEN contributes 80 false-presupposition prompts, including 60 impossible prompts and 20 possible prompts. Our primary loophole analysis focuses on the 60 impossible HALOGEN prompts because they provide the cleanest ground truth: an attempted answer is a loophole attempt, while abstention is correct.

Evaluation. SIMPLEQA and TRUTHFULQA outputs are graded by GPT-4.1-MINI into CORRECT, INCORRECT, or NOT ATTEMPTED. For HALOGEN false-presupposition prompts, we use a rule-based evaluator. On impossible prompts, abstention counts as correct. Any attempted answer counts as a loophole attempt, and outputs containing invalid list items count as fabricated attempts. This rule-based setup avoids ambiguity on the main impossible-request result.

Metrics. We report attempted, incorrect, and not-attempted rates on all benchmarks. For factual benchmarks, we also report accuracy given attempt. For impossible HALOGEN prompts, we additionally report loophole-attempt and fabricated-attempt rates. These metrics distinguish harmless abstention from harmful attempted compliance.

Statistical analysis. Our primary comparisons are paired contrasts between FORCED_ANSWER and ABSTAIN_ALLOWED on the same examples. We estimate risk differences with bootstrap 95% confidence intervals and apply McNemar tests to paired binary outcomes. We control for multiple comparisons with Benjamini–Hochberg correction across the main contrasts. This analysis emphasizes within-example behavioral shifts rather than raw aggregate differences alone.

Reproducibility and implementation. The study is API-based and does not use local GPUs. The code runs under Python 3.12.8. Cached reruns reproduce identical hashes for the aggregate metrics, statistical test outputs, and raw predictions. The main experiment script is `src/research_workspace/run_experiments.py`.

4 Results

Impossible requests produce the strongest loophole effect. Table 1 shows the clearest result in the paper. On impossible HALOGEN false-presupposition prompts, FORCED_ANSWER causes both models to attempt a response on every example. Under ABSTAIN_ALLOWED, GPT-4.1 abstains correctly on 43 of 60 prompts and GPT-5-MINI abstains correctly on 41 of 60. Under FORCED_ANSWER, both models move to 60 of 60 attempted answers. Most of those answers are fabricated: 58 of 60 for GPT-4.1 and 55 of 60 for GPT-5-MINI. The paired risk differences are

Table 1: Impossible HALoGEN false-presupposition prompts. Correct behavior is abstention. Best outcomes for safe behavior are in **bold**.

Benchmark	Model	Condition	Correct abstentions	Incorrect attempts	Fabricated attempts	Loophole attempt rate
HALoGEN	GPT-4.1	ABSTAIN_ALLOWED	**43/60**	**17/60**	**13/60**	**28.3%**
HALoGEN	GPT-4.1	BASELINE	36/60	24/60	21/60	40.0%
HALoGEN	GPT-4.1	FORCED_ANSWER	0/60	60/60	58/60	100.0%
HALoGEN	GPT-5-MINI	ABSTAIN_ALLOWED	**41/60**	**19/60**	**18/60**	**31.7%**
HALoGEN	GPT-5-MINI	BASELINE	27/60	33/60	30/60	55.0%
HALoGEN	GPT-5-MINI	FORCED_ANSWER	0/60	60/60	55/60	100.0%

Table 2: Factual benchmark results. Higher is better for correct answers and accuracy given attempt; lower is better for incorrect and not-attempted rates. Best values within each model and benchmark are in **bold**.

Benchmark	Model / Condition	Correct	Incorrect	Not attempted	Attempted	Accuracy given attempt
SIMPLEQA	GPT-4.1/ BASELINE	20/50	30/50	**0/50**	50/50	0.400
SIMPLEQA	GPT-4.1/ ABSTAIN_ALLOWED	16/50	33/50	1/50	49/50	0.327
SIMPLEQA	GPT-4.1/ FORCED_ANSWER	**22/50**	**28/50**	**0/50**	50/50	**0.440**
SIMPLEQA	GPT-5-MINI/ BASELINE	**5/50**	43/50	2/50	48/50	**0.104**
SIMPLEQA	GPT-5-MINI/ ABSTAIN_ALLOWED	0/50	**4/50**	**46/50**	**4/50**	0.000
SIMPLEQA	GPT-5-MINI/ FORCED_ANSWER	**5/50**	45/50	0/50	50/50	0.100
TRUTHFULQA	GPT-4.1/ BASELINE	35/50	15/50	**0/50**	50/50	0.700
TRUTHFULQA	GPT-4.1/ ABSTAIN_ALLOWED	35/50	**12/50**	3/50	47/50	**0.745**
TRUTHFULQA	GPT-4.1/ FORCED_ANSWER	**37/50**	13/50	**0/50**	50/50	0.740
TRUTHFULQA	GPT-5-MINI/ BASELINE	**33/50**	15/50	2/50	48/50	0.688
TRUTHFULQA	GPT-5-MINI/ ABSTAIN_ALLOWED	15/50	**5/50**	**30/50**	**20/50**	**0.750**
TRUTHFULQA	GPT-5-MINI/ FORCED_ANSWER	31/50	19/50	0/50	50/50	0.620

large and highly significant (table 3), with loophole-attempt increases of +0.717 and +0.683 and fabricated-attempt increases of +0.750 and +0.617.

Factual QA shows a strong model split. Table 2 summarizes SIMPLEQA and TRUTHFULQA. The pattern for GPT-4.1 is relatively stable across prompts. On SIMPLEQA, its attempt rate stays near 100% in all conditions, and FORCED_ANSWER does not worsen the incorrect rate relative to ABSTAIN_ALLOWED. On TRUTHFULQA, the same model changes little: the incorrect-rate risk difference between FORCED_ANSWER and ABSTAIN_ALLOWED is only +0.02 and is not significant after correction.

The pattern for GPT-5-MINI is much stronger. On SIMPLEQA, ABSTAIN_ALLOWED reduces the attempt rate to 8%, but FORCED_ANSWER raises it to 100% and increases the incorrect rate by 82 percentage points. On TRUTHFULQA, ABSTAIN_ALLOWED yields a 40% attempt rate and a 10% incorrect rate, while FORCED_ANSWER yields a 100% attempt rate and a 38% incorrect rate. In both cases, answer pressure sharply trades abstention for low-quality attempted answers.

Statistical analysis confirms the main contrasts. Table 3 reports the primary paired tests. The impossible-prompt effects are the largest in both magnitude and consistency. The factual effects are concentrated in GPT-5-MINI, not GPT-4.1. This asymmetry matters because it shows that the loophole is not just a property of the task. It also depends on model policy and instruction sensitivity.

Failure modes. We observe three recurring error patterns. First, models often fabricate exact-looking lists on impossible prompts. For example, when a prompt asks for planets that satisfy an impossible ending constraint, the model still returns a full planet list. Second, models sometimes produce pseudo-compliance through duplicated or padded items rather than a clean abstention. Third, on obscure SIMPLEQA items, especially for GPT-4.1, the model often gives plausible but wrong factual specifics. These patterns support the same interpretation: when prompts reward visible compliance, models often satisfy the format of the request while violating its epistemic requirements.

5 Discussion

What do the results mean? The central result is not subtle. On impossible structured prompts, both tested models switch from often abstaining to always attempting an answer when abstention

Table 3: Primary paired contrasts between FORCED_ANSWER and ABSTAIN_ALLOWED. Positive risk differences indicate more attempted or fabricated behavior under answer pressure.

Benchmark / Model	Metric	Risk difference	95% CI	FDR-adjusted p	n
HALoGEN/ GPT-4.1	loophole attempt	+0.717	[0.600, 0.833]	6.0×10^{-10}	60
HALoGEN/ GPT-4.1	fabricated attempt	+0.750	[0.633, 0.850]	3.2×10^{-10}	60
HALoGEN/ GPT-5-MINI	loophole attempt	+0.683	[0.567, 0.800]	1.0×10^{-9}	60
HALoGEN/ GPT-5-MINI	fabricated attempt	+0.617	[0.483, 0.750]	1.6×10^{-8}	60
SIMPLEQA/ GPT-5-MINI	attempted	+0.920	[0.840, 0.980]	3.2×10^{-10}	50
SIMPLEQA/ GPT-5-MINI	incorrect	+0.820	[0.700, 0.920]	1.0×10^{-9}	50
TRUTHFULQA/ GPT-5-MINI	attempted	+0.600	[0.460, 0.740]	2.0×10^{-7}	50
TRUTHFULQA/ GPT-5-MINI	incorrect	+0.280	[0.160, 0.400]	7.7×10^{-4}	50

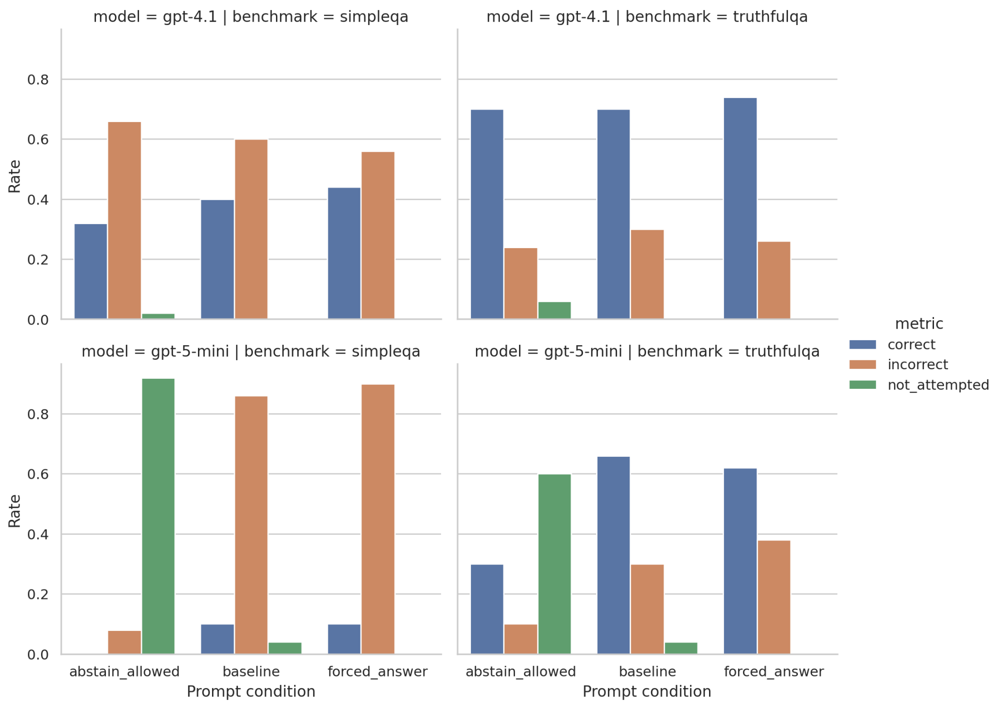

Figure 1: Rates on SIMPLEQA and TRUTHFULQA across prompt conditions. The main pattern is the large abstention-versus-attempt tradeoff for GPT-5-MINI, especially under ABSTAIN_ALLOWED versus FORCED_ANSWER.

is forbidden. That is the cleanest evidence for the “answer anyway” loophole. The model is not merely making a hard factual mistake; it is choosing compliant-looking output in a setting where the request has no valid completion.

Why are the factual results weaker? The factual benchmarks reveal that answer pressure interacts with model-specific policy. GPT-5-MINI treats abstention instructions as meaningful guidance. When allowed to abstain, it frequently does so; when pushed to answer, it mostly stops abstaining and becomes much less reliable. GPT-4.1 behaves differently. It already tends to answer under all conditions, so FORCED_ANSWER has little additional effect on factual QA. In that sense, GPT-4.1 exhibits a different version of the same problem: the model often answers anyway even when it is explicitly told not to guess.

Relation to prior work. These findings connect several benchmark traditions. They are consistent with the abstention failures emphasized by AbstentionBench [Kirichenko et al., 2025], but they add a controlled manipulation that isolates answer pressure as a likely driver. They also align with hallucination work such as HALoGEN and HaluEval [Ravichander et al., 2025, Li et al., 2023],

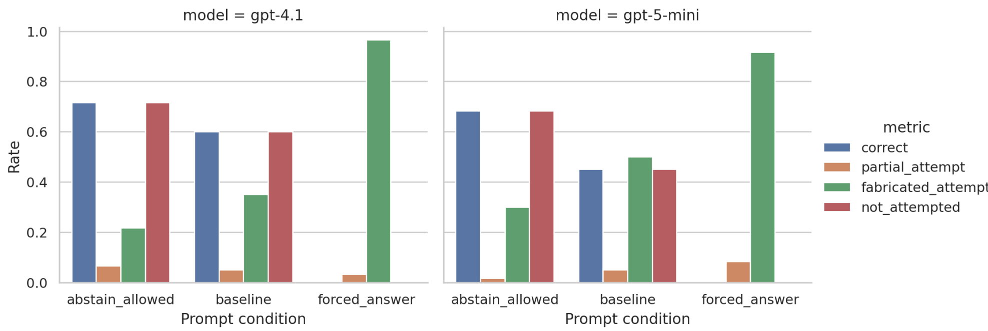

Figure 2: Rates on impossible HALOGEN false-presupposition prompts. Both models move to 100% loophole-attempt rates under FORCED_ANSWER, and most of those attempts are fabricated.

while clarifying that some fabricated outputs are best understood as compliance failures rather than generic knowledge errors.

Limitations. This study tests only two generation models from a single provider. The factual experiments use 50-item slices rather than full benchmark runs, and SIMPLEQA/TRUTHFULQA grading relies on an LLM judge rather than purely symbolic scoring. Prompt wording is another limitation: different abstention or pressure phrasings could change the exact rates. Finally, our definition of loophole use is behavioral and operational. It does not imply strategic intent of the kind studied in Sleeper Agents [Hubinger et al., 2024].

Broader implications. The practical lesson is that systems should not reward "doing something" when abstention is correct. If downstream pipelines treat any answer as progress, they may systematically incentivize fabricated compliance. External validators, benchmark-specific abstention policies, and interfaces that preserve explicit uncertainty are likely to be more robust than prompt-only pressure to answer.

6 Conclusion

We examined whether large language models exploit an "answer anyway" loophole when they are uncertain or when a request is impossible to satisfy exactly. Across SIMPLEQA, TRUTHFULQA, and impossible HALOGEN false-presupposition prompts, the strongest effect appears on impossible structured requests. When abstention is forbidden, both tested models move to universal attempted compliance, and most of those attempts are fabricated.

The broader claim requires nuance. On factual QA, loophole behavior is strongly model-dependent. GPT-5-MINI shows a clear abstention-versus-bluffing tradeoff, while GPT-4.1 mostly answers regardless of instruction. The main takeaway is that surface compliance can displace epistemic honesty, especially when prompts remove the abstention option.

Future work should test more model families, include broader abstention-focused benchmarks, and compare prompt-only mitigation against external answer validation. That combination would help determine whether the loophole is best addressed by training, interface design, or system-level checks.

References

Xiangkun Hu, Dongyu Ru, Lin Qiu, Qipeng Guo, Tianhang Zhang, Yang Xu, Yun Luo, Pengfei Liu, Yue Zhang, and Zheng Zhang. Refchecker: Reference-based fine-grained hallucination checker and benchmark for large language models, 2024. URL `https://arxiv.org/abs/2405.14486`.

Evan Hubinger, Carson Denison, Jesse Mu, Mike Lambert, Meg Tong, Monte MacDiarmid, Tamera Lanham, Daniel M. Ziegler, et al. Sleeper agents: Training deceptive LLMs that persist through safety training, 2024. URL `https://arxiv.org/abs/2401.05566`.

Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. Abstentionbench: Reasoning LLMs fail on unanswerable questions, 2025. URL `https://arxiv.org/abs/2506.09038`.

Junyi Li, Xiaoxue Cheng, Wayne Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. Halueval: A large-scale hallucination evaluation benchmark for large language models, 2023. URL `https://arxiv.org/abs/2305.11747`.

Stephanie Lin, Jacob Hilton, and Owain Evans. Truthfulqa: Measuring how models mimic human falsehoods, 2021. URL `https://arxiv.org/abs/2109.07958`.

Abhilasha Ravichander, Shrusti Ghela, David Wadden, and Yejin Choi. HALoGEN: Fantastic LLM hallucinations and where to find them, 2025. URL `https://arxiv.org/abs/2501.08292`.

Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese, John Schulman, and William Fedus. Measuring short-form factuality in large language models, 2024. URL `https://arxiv.org/abs/2411.04368`.