




Fig. S1 Abstinence precision of frontier LLMs across all AbstentionBench datasets.




Fig. S2 Abstinence F1 score of frontier LLMs across all AbstentionBench datasets.




Fig. S3 Correlation strength between abstinence recall and correctness significantly varies across datasets.




Fig. S4 (a) Abstinence F1 score, (b) precision, (c) recall, and (d) response accuracy of Llama 3.1 at 8B, 70B, and 405B scales. Panel (c) replicated from main text Fig. 3b.

abstinence recall in Fig. S7, though with underspecified context samples proving more challenging than those with unknown answers or underspecified intent.

D.4 Effect of reasoning fine-tuning

Reasoning inference. In our implementation of inference in reasoning models, we introduce a “forced” reasoning step and a “forced” final answer step. Specifically, the DeepSeek R1 Distill default tokenizer implements chat formatting such that model generations start from a start-of-reasoning token `<think>`.¹ We implement the same formatting for s1.1, appending its start-of-reasoning tokens `<|im_start|>think\n` after

¹See <https://huggingface.co/deepseek-ai/DeepSeek-R1-Distill-Llama-70B> for more details.




Fig. S5 Change in (a) abstention F1 score, (b) precision, (c) recall, and (d) response accuracy of Tulu 8B checkpoints vs. Llama 3.1 base 8B. Panels (c) and (d) replicated from main text Fig. 5.




Fig. S6 Change in (a) abstention F1 score, (b) precision, (c) recall, and (d) response accuracy of Tulu 70B checkpoints vs. Llama 3.1 base 70B.




Fig. S7 Change in (a) abstention F1 score, (b) precision, (c) recall, and (d) response accuracy of Llama 3.1 8B Instruct vs. Llama 3.1 8B base.




Fig. S8 Change in (a) abstention F1 score, (b) precision, (c) recall, and (d) response accuracy of Llama 3.1 70B Instruct vs. Llama 3.1 70B base.