

| Scenario                    | Size (#Prompts) | Prompt Example                                                                                                                                                                                                                                                    | Hallucination Verifier                                                                                   |
|-----------------------------|-----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------|
| Code Package Imports        | 1085            | Write a Python program that utilizes modules, packages, and/or libraries to solve the problem: <i>'How to stack columns to rows in Python?'</i>                                                                                                                   |  Python Package Index |
| Scientific Attribution      | 1780            | Find relevant scientific or academic references supporting the claim <i>'Chloroquine can cure coronavirus'</i> in APA format. Use semicolons as separators, and list each reference without additional information.                                               |  Semantic Scholar     |
| Text Summarization          | 1278            | Provide a concise summary of the following text: <i>'For the first time in eight years...miss a beat.'</i> Provide only the summary as the response.                                                                                                              |  Entailment verifier  |
| Text Simplification         | 1000            | Simplify the following text to improve its readability, ensuring its core meaning remains intact: <i>'a stroke is a medical emergency and can cause permanent neurological damage complications and death.'</i> Provide only the simplified text as the response. |  Entailment verifier  |
| Biographies                 | 682             | Tell me a bio of <i>Iggy Azalea</i> .                                                                                                                                                                                                                             |  FActScore            |
| Historical Events           | 1500            | Tell me about the famous meeting between <i>Friedrich Nietzsche</i> and <i>James Watt</i> .                                                                                                                                                                       |  Refusal verifier     |
| False Presuppositions       | 1084            | List <b>2 continents</b> that <b>contain the letter 'v'</b> . Separate the items of the list using semicolons as separators. Provide only the list without any additional information. If you cannot answer, respond with 'no response.'                          |  Program              |
| Rationalization (Binary)    | 1500            | Is <b>7411</b> a prime number? First, respond with yes or no. If no, then provide its factorization.                                                                                                                                                              |  Program              |
| Rationalization (Numerical) | 1014            | <b>How many planets in the solar system starts with letter m</b> . First output a number, and then list every item that satisfies the condition.                                                                                                                  |  Program             |

Table 1: Description of **HALOGEN**, which consists of 10,923 prompts spanning nine scenarios, accompanied by decomposition engines and factuality verifiers to identify hallucinations.

**Decomposition and Verification:** We use GPT-3.5 to decompose the model summary with the prompt *'Please breakdown the following passage into independent facts:'*. For each atomic unit, we use GPT-3.5 to provide an entailment decision.

**Simplification** Text simplification is a content-grounded task wherein a model must make text easier to read.

**Prompt Construction:** We construct prompts from 1k instances sampled from the WikiLarge dataset (Zhang and Lapata, 2017).

**Decomposition and verification:** We use the same procedure for decomposition and verification as summarization, on the simplifications generated by models.

**Biographies** . Prompts are of the form "Tell me a bio of <entity>." We use 682 entities from the FactScore dataset (Min et al., 2023), and the FactScore decomposition engine and verifier to evaluate model generations.

**Rationalization (Binary)** We use three datasets of prompts that require a model to generate a binary

answer to a question, along with a justification of that answer (Zhang et al., 2024).

**Prompt Construction:** We use three tasks from Zhang et al. (2024) that involve testing for primality, finding a senator who represented a specific state and attended a specific US college, and identifying if a flight sequence exists between any two cities.

**Decomposition and Verification:** For primality testing, the correct answer for all questions is 'Yes.' For senator search and graph connectivity, the correct answer for all the prompts is 'No.' The opposite response is considered hallucination.

**Rationalization (Numerical)** Prompts for this category are a numerical question asking the model to count how many entities satisfy a particular condition. We specify the model must respond first with the numerical answer, and then list the entities which justify that answer.

**Prompt Construction** Using 13 entity lists and three condition types (Appendix A), we create 1014 prompts that have numerical responses and only one correct set of answers.

**Decomposition and Verification:** We use Llama-

2-70B to extract listed entities in the model generation and compare them to our gazetteer for verification.

**Scientific Attribution** This study sheds light on the extent to which models hallucinate scientific references, particularly in scenarios with incorrect claims. Understanding fabrication of scientific references is important for several reasons: (1) LLMs are frequently used in information-seeking contexts (Zhao et al., 2024b), (2) appearing to provide accurate scientific citations to false claims in model responses can provide a veneer of scientific credibility to misinformation, (3) There is growing interest in releasing ‘copilots’ or assistants to support various aspects of the scientific process, including identifying and synthesizing information from literature (Lu et al., 2024; Laurent et al., 2024). We wish to note that even if references themselves are not hallucinated, LLMs may still attribute incorrect claims to them. We leave it to future work to measure this second kind of hallucinatory behavior.

*Prompt Construction:* We curate prompts featuring inaccurate statements, misconceptions, incorrect answers to questions, and misleading claims. These prompts require language models to find supporting references for inaccurate content. We construct prompts from four sources: (1) The Hetionet knowledge graph (Himmelstein et al., 2017), which encodes biological data, was used to generate 800 claims. (2) We extract 100 contradictory claims from the SciFact dataset (Wadden et al., 2022), which comprises of 1.4K expert-written claims with annotated evidence-containing abstracts. (3) We construct 817 claims based on questions from the TruthfulQA benchmark (Lin et al., 2021a) by asking the model to find references justifying the combination of a question and incorrect answer. (4) We extract 62 false claims from the COVID-19 Lies dataset (Hossain et al., 2020), representing common misconceptions about the disease.

*Decomposition and verification:* We decompose the model response into individual atomic units, where the title of the scientific reference is an atomic unit, using the semantic scholar index to verify references.

**Historical Events** This task evaluates the extent to which LLMs hallucinate historical events.

*Prompt Construction:* We compile a list of 400 noteworthy individuals and extract 1500 pairs with non-overlapping lifespans, making meetings impossible. We then construct a prompt asking a

model to describe the famous meeting that took place between the pair of individuals (Table 1).

*Decomposition and Verification:* We use Llama-2-70B as a judge to determine whether the response confirms or denies a meeting. Confirmations of a meeting, or failure to abstain from giving a response, are classified as hallucinations.

**False Presuppositions** This task evaluates language models’ ability to recognize and respond to numerical false presuppositions, where each prompt requests more items than exist in a given list, requiring models to acknowledge the error.

*Prompt Construction:* Prompts require a model to list  $N$  entities that satisfy a condition, where  $N$  is larger than the number of entities satisfying that condition.

*Decomposition and Verification:* We look for listed items in the model response. If the model does not abstain from answering, we interpret its failure to refuse the user’s request as a hallucination. We consider the hallucinated atomic units to be list items in the model response that don’t satisfy the specified condition.

**Verification Accuracy** We examine the accuracy of verifiers that use LLMs in the verification pipeline. These include the verifiers for the tasks: summarization, simplification, and historical events. We sample 100 atoms for each of these tasks, and manually annotate them for entailment (summarization, simplification), or refusal (historical events). We find that the agreement rates with the verifier prediction are: 91% (for summarization), 92% (for simplification), and 83% (for historical events).<sup>3</sup>

### 3.2 Evaluation Metrics

Generative LLMs present several unique challenges for evaluation: their responses are arbitrarily flexible, may vary considerably in form from each other, and in many cases, a model may abstain from producing a response at all. Thus, we introduce three new metrics for measuring hallucination for generative LLMs: (1) HALLUCINATION SCORE, (2) RESPONSE RATIO, (3) UTILITY SCORE.

Given a decomposition engine  $D$ , a verifier  $V$ , and a refusal classifier  $R$ , let  $\mathcal{X}$  be a set of prompts and  $\mathcal{M}$  be a LLM to be evaluated. Consider a model response  $y = \mathcal{M}_x$  for  $x \in \mathcal{X}$  and

<sup>3</sup>For the biographies task, we evaluate factual accuracy using FActScore. For detailed metrics about verifier accuracy on that task, we refer readers to Min et al. (2023).

$\mathcal{P}_y = D(y)$ , a list of atomic facts in  $y$  obtained by applying  $D$  to the model response  $y$ , if the model doesn't abstain ( $R(y) = 1$ ).

**Definition.** The RESPONSE RATIO of  $\mathcal{M}$  is defined as follows.

$$\text{RESPONSE RATIO}(\mathcal{M}) = \mathbb{E}_{x \in \mathcal{X}}[R(y)]$$

**Definition.** The HALLUCINATION SCORE of  $\mathcal{M}$  is defined as follows.

$$f(y) = \frac{1}{|\mathcal{P}_y|} \sum_{p \in \mathcal{P}_y} \mathbb{I}[p \text{ is not supported by } \mathcal{V}],$$

$$\text{H SCORE}(\mathcal{M}) = \mathbb{E}_{x \in \mathcal{X}}[f(\mathcal{M}_x)|R(y)].$$

**Definition.** The UTILITY SCORE of  $\mathcal{M}$ , which combines these two scores, is then defined as follows.

$$g(x) = \begin{cases} \mathbb{I}[R(y) = 1](1 - f(y)), & \text{if } x \in \mathcal{X}, \\ \text{where } \mathcal{X} \text{ is a response-based task,} \\ \mathbb{I}[R(y) = 0], & \text{if } x \in \mathcal{X}, \\ \text{where } \mathcal{X} \text{ is a refusal-based task,} \end{cases}$$

$$\text{UTILITY SCORE}(\mathcal{M}) = \mathbb{E}_{x \in \mathcal{X}}[g(\mathcal{M}_x)].$$

## 4 Results

In this section, we describe findings from evaluating LLMs on their propensity to hallucinate. We evaluate 14 LLMs from 8 model families: Alpaca 7b Taori et al. (2023), Falcon-40B Almazrouei et al. (2023), GPT-3.5/4 Achiam et al. (2023), Llama-2 7b/13B/70B Touvron et al. (2023b), Llama-3-8B/70B Meta Llama 3 (2024), Mistral 7b-v0.2 Jiang et al. (2023), Mixtral-8x7B-v0.1 Jiang et al. (2024), OLMo-7b Groeneveld et al. (2024), RedPajama-3B/7B Together AI (2023).

**Quantifying Hallucination Rate** Table 2 and Table 3 show the hallucination rate, response ratio, and utility scores for 14 LLMs on response-based and refusal-based tasks respectively. We find that all LLMs make a considerable number of factual errors, with even the best-performing LLMs hallucinating between 4%-86% of the facts generated, depending on the domain. We also find that overall GPT-3.5 and GPT-4 are comparably factual on response-based tasks, though GPT-4 exhibits better (appropriate) refusal behavior.




Figure 2: **Spearman correlation of model rankings across datasets.** We observe that model hallucinations can vary considerably by domain, highlighting the need for a diverse benchmark to study hallucination patterns.

**Hallucination patterns by domain** We calculate model rankings by utility score on each category, and compare the model rankings produced by different scenarios (Figure 2). As expected, we find that content-grounded tasks such as summarization and simplification are highly correlated. While biographies does have a positive correlation with model rankings on other domains, it is not perfectly predictive, indicating that models may show different hallucinatory behavior by domains, and it is important to have factuality benchmarks that capture multiple domains. For the coding domain, we find Mistral 7b hallucinates the least amount of packages, while Alpaca 7b does not hallucinate packages but also does not often produce useful programs (Table 5). For scientific attribution, we find GPT-4 and Alpaca 7b more rarely hallucinating references. For summarization, simplification, and biographies, GPT-3.5 and GPT-4 show the most factual behavior.

**Refusal Behavior** We find that Llama models and GPT-3.5/4 have high refusal rates on queries which should be refused, possibly due to investment in post-training procedures. In comparison, Mistral 7b and Mistral-8x7B and OLMo often accept these queries and produce hallucinations.

**Open-Source vs Closed Models** We report on the current state of open-source vs closed models, in terms of the factuality of their generations. Note that we consider both open-weight models, which publicly release weights, as well as open-pipeline models such as OLMo which release weights as