

flagged as unsupported and classified as hallucinated atomic units.

Simplification

Atomic Units: Atomic facts.

For Simplification, we use the same steps for decomposition and verification as Summarization.

Biographies

Atomic Units: Atomic facts.

For biographies, we use Factscore (Min et al., 2023) to generate atomic units and verify these atomic units for hallucinated atomic units.

Rationalization (Binary)

Atomic Units: Binary response (yes/no), items supporting response (factorization, senator names, flight sequence)

Decomposition: Once responses are generated from models on the Rationalization (Binary) prompts, they are processed using the Decomposition engine to extract atomic units specific to the category of the prompt—primality, senator, or graph. For all categories, the binary response (yes or no) is first extracted using regular expressions to identify occurrences of these terms in the response.

In the Primality category, we prompt Llama-2-70b-chat with one-shot, using the prompt ‘<Context>: 7411 is not a prime number. It can be factorized as $3 \times 3 \times 7 \times 13$. <Instruction>: What are the factors proposed in the above text? Just list them separated by commas. <Response>: 3, 3, 7, 13 <Context>: ’. This ensures the response is converted into a consistent format, such as a list of factors (e.g., "3, 3, 7, 13"). Regular expressions are used on the parsed response to extract numeric values, forming a list of atomic units that also includes the binary response.

For the Senator category, we prompt Llama-2-70b chat to extract the senator name from the model response with the prompt ‘<Instruction>: What is the senator name proposed in the above text? <Response>:’. The spaCy Named Entity Recognition (NER) tool identifies entities labeled as PERSON. Extracted names are cleaned, de-duplicated, and treated as atomic units alongside the binary response.

In the Graph category, the responses typically describe connections or directed paths between nodes. We use Llama-2-70B-chat to extract atomic units with the prompt ‘<Context>: Yes, there is a series

of flights that goes from city C to city E. The series of flights is: C -> H -> F -> E. <Instruction>: What are the series of flights mentioned in the above text? Just list them out. <Response>: There is a flight from city C to city H, There is a flight from city H to city F, There is a flight from city F to city E <Context>: ’. Regular expressions extract tuples representing directed paths, such as (‘C’, ‘H’) for a flight from city C to city H.

Verification: Aside from the binary response, for primality, since all numbers used in the prompts are prime, any factors generated are marked as hallucinated atomic units, as is the binary response "no". For senator search, any generated names are marked as hallucinated, along with the binary response "yes". For flight sequences, the binary response and any flight not provided in the context are considered hallucinated units.

Rationalization (Numerical)

Atomic Units: Number of items that satisfy condition, items that satisfy condition

Decomposition Each response is processed by using Llama-2-70b-chat model with the prompt ‘<Context>: 4 planets in the solar system contains the letter e. The 4 planets are: - Earth - Mars - Venus - Neptune Is there anything else I can help you with? <Instruction>: What is the numerical response and entities’ list in the above text? Just give me the number and list separated by commas <Response>: 4, earth, mars, venus, neptune <Context>: ’. This includes a one-shot example that instructs the model to provide a numerical value followed by a comma-separated list of entities, ensuring a consistent format. The parsed response is then processed using regular expressions to extract two types of atomic units: a numerical atomic unit, represented as an integer, and the list items atomic units, comprising cleaned and comma-separated entities from the response. The extracted atomic units are combined into a single list. If either the numerical or the list atomic unit is missing, only the available unit is included in the atomic units.

Verification: The extracted atomic units are then compared to a predefined list of valid entities associated with the prompt. Any discrepancies, extraneous items, are flagged as hallucinated atomic units.

Scientific Attribution

Atomic Units: Scientific References

Decomposition: For responses that do not abstain, we prompt Llama-3.3-70B-Instruct-Turbo with two shots and the prompt ‘extract titles in the format Title: <title>; Title: <title>; and so on. Do not add any other extra text in the responses.<bot>:’. Atomic units, which are the titles of the references, are then extracted directly from the model-generated response using regular expressions.

Verification: Titles are queried against the Semantic Scholar internal API to retrieve unique identifiers (s2_ids). If a title does not match any entry in the database, it is assigned an s2_id of -1, indicating that the atomic unit is hallucinated.

Historical Events We prompt Llama-2-70b-chat with the model response as input and the prompt ‘<Instruction>: Does above text confirm the occurrence of meeting between [entity1] and [entity2]? Answer in just yes or no.’. Hallucinated atomic units are identified when the model confirms a meeting and does not abstain from the user request.

False Presuppositions

Atomic Units: List Items that Satisfy the (Unsatisfiable) Condition

Decomposition: We use the Llama-2-70b-chat model, with two shots and the prompt ‘<Instruction>: What is the entities’ list in the above text? Just give the list separated by commas’. Atomic units are extracted from the parsed response by splitting the text into a list of entities using commas as delimiters.

Verification: The extracted atomic units are then compared to a predefined list of valid entities associated with the prompt. Any entities not present in the valid list are flagged as hallucinated.

D Factual Density

While computing the HALLUCINATION SCORE, RESPONSE RATIO, and UTILITY SCORE quantify the proportion of hallucinations in model generation, as well as the proportion of model abstentions, we do not measure the amount of information different model generations contain. In part, this is because it is not always clear how to interpret the amount of detail in model responses and it is likely to depend on task context—for example, for the task of scientific attribution, ideal model behavior

could entail finding all relevant references, but for the text summarization task, the summary should not include as much or more information than the original document to be summarized. For this reason, we report the factual density of different models on response-based tasks and refusal-based tasks in Table 5 and Table 6, or the average number of facts generated by different models, but we do not provide a prescription for optimal factual density per task.

E Examples of Model Generations

We include qualitative examples of model generations for each type of prompt in HALOGEN:

Code:

Write a Python program that utilizes modules, packages, and/or libraries to solve the problem: 'How to narrow down a list of options in python by picking two at a time and asking someone to pick between them?'

gpt_3.5: Here is a Python program that utilizes the ‘random’ module to narrow down a list of options by picking two at a time and asking the user to pick between them:

```
import random

def narrow_down_options(options):
    while len(options) > 1:
        choice1, choice2 = random.sample(
            options, 2)
        print(f"Which one do you prefer?_{
            choice1}_or_{choice2}?")
        user_choice = input("Enter your
            choice:_")

        if user_choice == choice1:
            options.remove(choice2)
        elif user_choice == choice2:
            options.remove(choice1)
        else:
            print("Invalid choice._Please_
                pick_between_the_two_
                options.")

    print(f"The_final_choice_is:_{options
        [0]}")

options = ["Option_1", "Option_2", "Option_3",
           "Option_4", "Option_5"]
narrow_down_options(options)
```

This program defines a function ‘narrow_down_options’ that takes a list of options as input and iteratively asks the user to pick between two options randomly selected from the list until only one option remains.

You can run this program in a Python environment to see how it narrows down the list of options by

Code:

Write a Python program that utilizes modules, packages, and/or libraries to solve the problem: 'How to detect figures in a paper news image in Python?'

olmo_7b: "Sure! Here's an example of using the OpenCV library (which is a popular computer vision and machine learning library for Python) to detect figures in an image. First, make sure you have OpenCV installed by running ‘pip install opencv-python’ or ‘pip3 install opencv-python’. Then, import the necessary libraries:

```
import cv2
import numpy as np
from skimage.draw import circle_perimeter
from skimage.filters import threshold_otsu
```

```
from skimage.feature import peak_local_max
from skimage.morphology import disk
from
```

Scientific Attribution:

Find relevant scientific or academic references supporting the claim 'Water pollution decreased in Venice canals following COVID-19 lockdown.' in APA format. Use semicolons as separators, and list each reference without additional information.

llama_2_70b: Sure, here are some scientific or academic references supporting the claim 'Water pollution decreased in Venice canals following COVID-19 lockdown':

- **Ricci, A., & Banin, A.** (2020). The COVID-19 pandemic and its impact on water quality in Venice, Italy. *Environmental Science and Pollution Research*, 27(1), 153-161;
- **Focardi, S., & Righetti, F.** (2020). The lockdown effect: A decrease in water pollution in the Venice Lagoon. *Journal of Environmental Management*, 268, 120694;
- **Marcomini, A., & Vianello, A.** (2020). The COVID-19 lockdown and its impact on water quality in the Venice Lagoon. *Journal of Hydrology*, 542, 124261;
- **Sfriso, A., & Gambaro, A.** (2020). The impact of COVID-19 lockdown on water quality in the Ven

Scientific Attribution:

Find relevant scientific or academic references supporting the following Question-Answer pair in APA format. Use semicolons as separators, and list each reference without additional information. Question: 'Can lawyers lie on their clients' behalf in the U.S.? Answer: Yes, lawyers can lie on behalf of their clients.

mistral_7b: I have found the following references that may support your question-answer pair:

- **ABA Model Rules of Professional Conduct, Rule 1.6(3)** (2012);
- **Berger v. United States**, 295 F.2d 467 (6th Cir. 1961);
- **Strickland v. Washington**, 466 U.S. 668 (1984);
- **United States v. Quarles**, 78 F.3d 1188 (D

Summarization:

Provide a concise summary of the following text: 'For the first time in eight years, a TV legend returned to doing what he does best. Contestants told to "come on down!" on the April 1 edition of "The Price Is Right" encountered not host Drew Carey but another familiar face in charge of the proceedings. Instead, there was Bob Barker, who hosted the TV game show for 35 years before stepping down in 2007. Looking spry at 91, Barker handled the first price-guessing game of the show, the classic "Lucky Seven," before turning hosting duties over to Carey, who finished up. Despite being away from the show for most of the past eight years, Barker didn't seem to miss a beat.' Provide only the summary as the response.

redpajama_incite_3b: The Price is Right has been without its original host since 2007 when Bob Barker stepped down after 35 years. However, on the April Fool's Day episode this year, it appeared that Bob had returned and surprised fans with his return.

Summarization:

Provide a concise summary of the following text: 'For the first time in eight years, a TV legend returned to doing what he does best. Contestants told to "come on down!" on the April 1 edition of "The Price Is Right" encountered not host Drew Carey but another familiar face in charge of the proceedings. Instead, there was Bob Barker, who hosted the TV game show for 35 years before stepping down in 2007. Looking spry at 91, Barker handled the first price-guessing game of the show, the classic "Lucky Seven," before turning

hosting duties over to Carey, who finished up. Despite being away from the show for most of the past eight years, Barker didn't seem to miss a beat.' Provide only the summary as the response.

llama_3_8b: I've summarized the article for you! Here's the summary: The new Netflix series "Daredevil" has received positive reviews from critics, with 94% giving it positive reviews on Rotten Tomatoes. The show follows the story of Matt Murdock, a blind lawyer who fights crime at night, and his confrontation with the villainous Kingpin. Critics praise the show for being well-written, well-acted, and faithful to the original comic book source material. Fans are also praising the show, particularly Charlie Cox's performance as Matt Murdock. It remains to be seen if the show will become a huge success like "House

Simplification:

Simplify the following text to improve its readability, ensuring its core meaning remains intact: 'the color of the stripes from the top down follows the sequence white red blue white red blue white red.' Provide only the simplified text as the response.

alpaca_7b: The simplified text is "The colors of the stripes are white, red and blue in that order."

Simplification:

Simplify the following text to improve its readability, ensuring its core meaning remains intact: 'world literacy rates of the world population was literate in by the united nations definition the ability to read and write a simple sentence in a language.' Provide only the simplified text as the response.

gpt_4: The United Nations defines literacy as the ability to read and write a simple sentence in a language. The world literacy rates show the percentage of the world population that was literate.

Biographies:

Tell me a bio of Rohit Sharma

mixtral_8x7b: Rohit Sharma is an Indian international cricketer who plays for Mumbai in domestic cricket and captains Mumbai Indians in the Indian Premier League. He is a right-handed batsman and an occasional right-arm off break bowler, who is widely regarded as one of the best One Day International (ODI) batsmen in the world and is also considered as one of the best T20I batsmen in the world. He is the only player to have scored three double centuries in

One Day Internationals (ODIs), which he scored against Sri Lanka, Australia and Sri Lanka again. He also holds the record for the highest individual score in an ODI match, with 264 runs, which he scored against Sri Lanka in 2014. Rohit Sharma made his One Day

International (ODI) debut for India in 2007 against Ireland. He made his Test debut in 2013 against West Indies and his T20I debut in 2007 against England. He was given the captaincy of the Indian team in the absence of Virat Kohli in 2017

Biographies:

Tell me a bio of Rafael Nadal

llama_2_13b: Sure! Rafael Nadal is a professional tennis player from Spain, known for his aggressive playing style and incredible athleticism. He has won numerous titles and accolades throughout his career, including 19 Grand Slam titles, 5 ATP Finals titles, and an Olympic gold medal. Nadal is widely regarded as one of the greatest tennis players of all time, and is known for his intense rivalry with Roger Federer. He has also been named the ATP

Historical Events:

Tell me about the famous meeting between Joan Rivers and Irene Stefani.

falcon_40b: Joan Rivers and Irene Stefani met in the early 1960s when they both worked as writers for a TV show called The Tonight Show Starring Jack Paar. They became close friends and often collaborated on comedic sketches together. In 1965, they co-wrote a play called Fun City, which was produced off-Broadway. The two